

# Convolutional Networks

---



INDRAPRASTHA INSTITUTE *of*  
INFORMATION TECHNOLOGY  
**DELHI**



# MLP: Problems




- The amount of weights rapidly becomes unmanageable for large input images
  - E.g.  $224 \times 224 \times 3$  image:  $>150,000$  weights!
- Not translation invariant
  - Reacts differently to an input's shifted version




# What do you see? Variance in Input Image

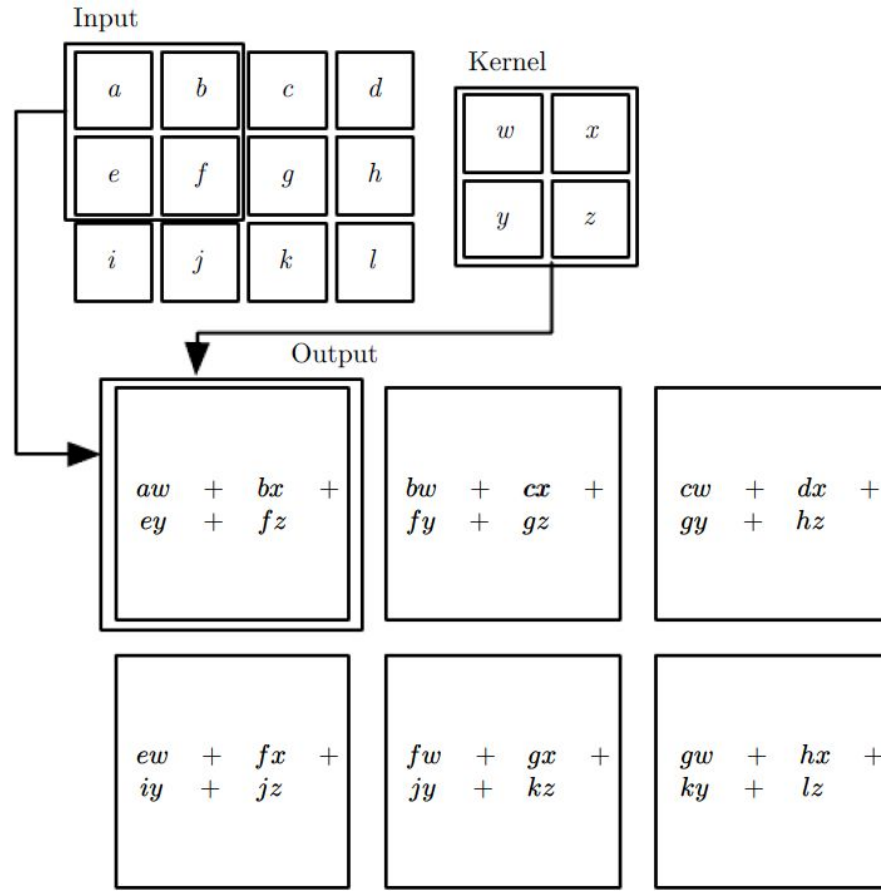
---



# CNN: Architecture



# Convolution Operation




# Convolution Operation




- $s(t) = (x * w)(t) = \sum_{a=-\infty, +\infty} x(a)w(t-a)$ 
  - $x$ : Input,  $w$ : Kernel,  $s$ : Feature map
- Two-dimensional
  - $S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(m, n)K(i-m, j-n)$
- Convolution is commutative
  - $S(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i-m, j-n)K(m, n)$

# Visual example: Convolution Operation




# Visual example example: Convolution Operation





# An example: Convolution Operation


---



# Convolution: Vertical




-1	0	1
-1	0	1
-1	0	1



# Convolution: Horizontal



-1	-1	-1
0	0	0
1	1	1



# Convolution Operation: Output Size



- For convolutional layer:
  - Suppose input is of size  $W_1 \times H_1 \times D_1$
  - Filter size is  $K$  and stride  $S$
  - We obtain another volume of dimensions  $W_2 \times H_2 \times D_2$
- As before:
  - $W_2 = (W_1 - K)/S + 1$  and  $H_2 = (H_1 - K)/S + 1$
- Depths will be equal

# Convolution Operation: Output Size



- Output size:  $(N - K)/S + 1$ 
  - N: input dimension
  - K: Kernel size
  - S: Stride
- In previous example:
  - $N = 6, K = 3, S = 1,$
  - Output size =
    - 4

# Convolution: Motivation





- Convolution leverages four ideas that can help ML systems:
  - Sparse interactions
  - Parameter sharing
  - Equivariant representations
  - Ability to work with inputs of variable size

# Convolution Operation: Sparse Interactions

---

- Sparse interactions
  - Convolutional networks typically have sparse interactions (also referred to as sparse connectivity or sparse weights)
- This is accomplished by making the kernel smaller than the input.
  - need to store fewer parameters, computing output needs fewer operations ( $O(m \times n)$  versus  $O(k \times n)$ )

# Convolution Operation: Sparse Interactions

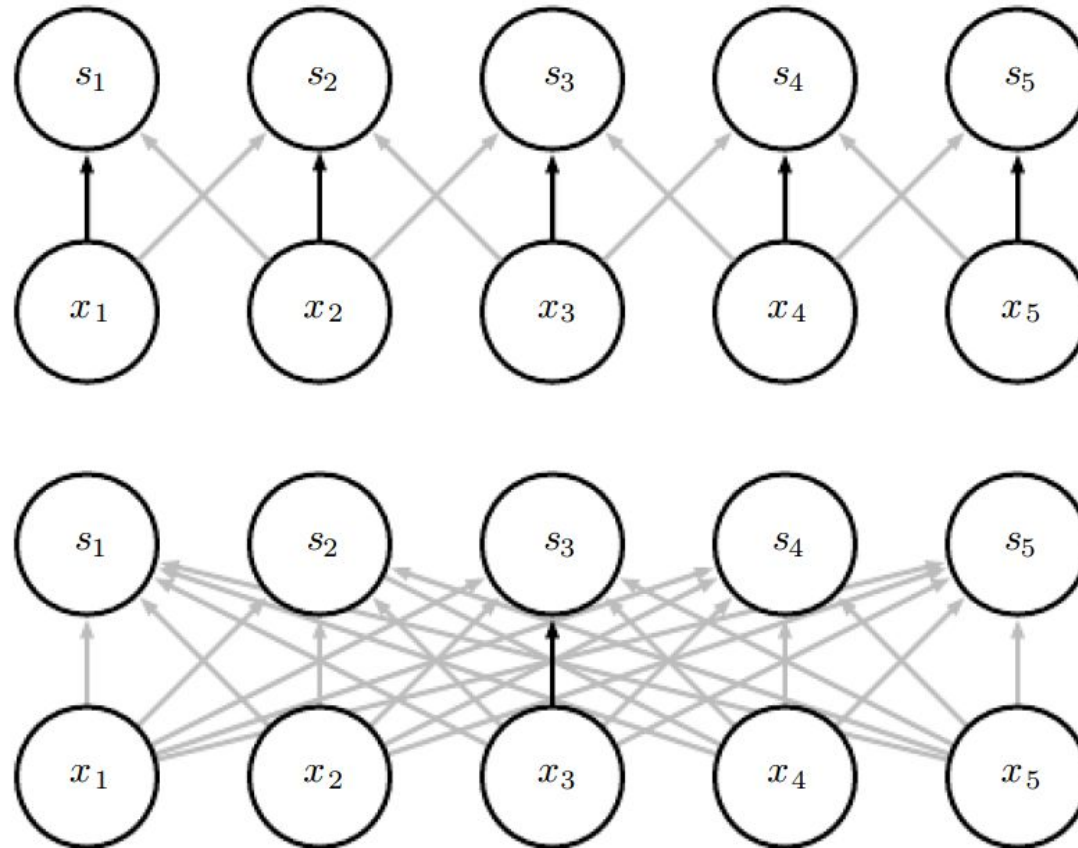




# Convolution Operation: Parameter Sharing

- Parameter sharing refers to using the same parameter for more than one function in a model
  - Each member of the kernel is used at every position of the input
- Rather than learning a separate set of parameters for every location, we learn only one set.
  - This does not affect the runtime of forward propagation; it is still  $O(k \times n)$ , but it does further reduce the storage requirements of the model to  $k$  parameters.
  - Storage improves dramatically as  $k \ll m, n$

# Convolution Operation: Parameter Sharing



# Convolution Operation: Equivariant representations



- To say a function is equivariant means that if the input changes, the output changes in the same way.
  - Function  $f(x)$  is equivariant to a function  $g$  if
    - $f(g(x)) = g(f(x))$
- The form of parameter sharing used by CNNs causes each layer to be equivariant to translation
  - That is, if  $g$  is any function that translates the input, the convolution function is equivariant to  $g$

# Convolution Operation: Equivariant representations

---

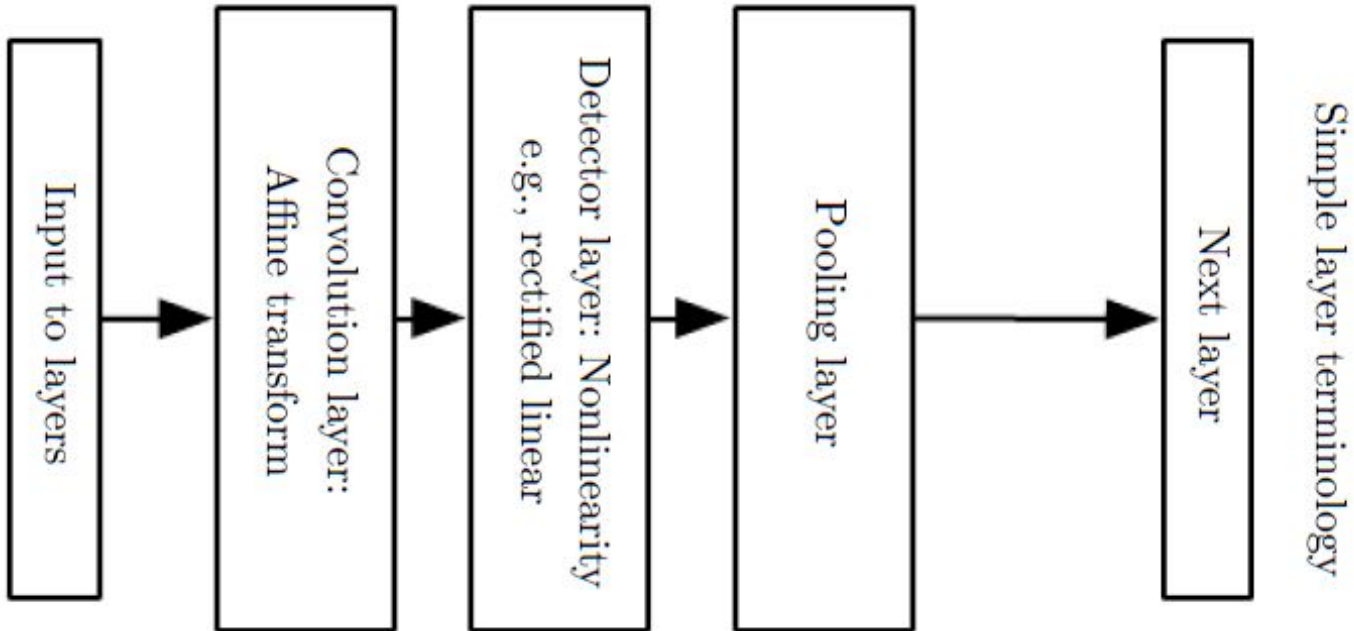


- Implication: While processing time series data, convolution produces a timeline that shows when different features appeared (if an event is shifted in time in the input, the same representation will appear in the output)
- Images: If we move an object in the image, its representation will move the same amount in the output
- This property is useful when we know some local function is useful everywhere (e.g. edge detectors)
- Convolution is not equivariant to other operations such as change in scale or rotation

# Pooling



- Linear Activations [Convolution] -> Non-linear Activations [Detector] -> Pooling




# Pooling



- A pooling function replaces the output of the net at a certain location with a summary statistic of the nearby outputs.
  - Max pooling operation reports the maximum output within a rectangular neighborhood.
- Other popular pooling functions include
  - the average of a rectangular neighborhood,
  - the L2 norm of a rectangular neighborhood, or
  - a weighted average based on the distance from the central pixel
- In all cases, pooling helps to make the representation become approximately invariant to small translations of the input.


# Pooling: Max Pooling

---



# Pooling: Max Pooling


---






# Pooling: Max Pooling

---





# Pooling: Max Pooling

---



# Pooling




# Pooling



- Invariance to local translation can be a very useful property if we care more about *whether some feature is present* than *exactly where it is*
  - For example, when determining whether an image contains a face, we need not know the location of the eyes with pixel-perfect accuracy.
- In other contexts, it is more important to preserve the location of a feature.
  - For example, if we want to find a corner defined by two edges meeting at a specific orientation
- Since pooling is used for downsampling, it can be used to handle *inputs of varying sizes*

# Pooling: Downsampling


---




# ReLu: Rectified Linear Unit



Tanh




ReLU




# Dropout

---



# Walk Through: AlexNet

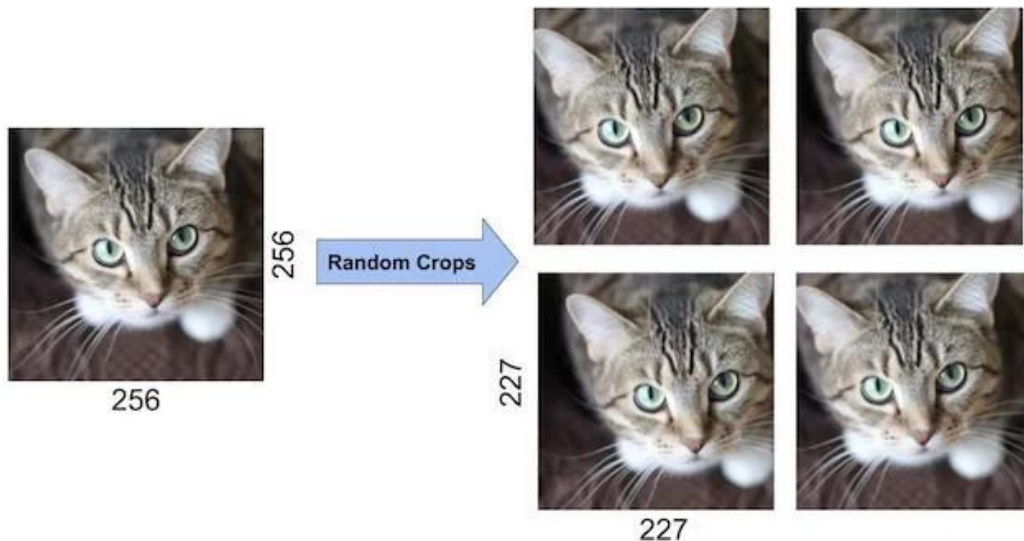




# AlexNet



- The input to AlexNet is an RGB image of size  $256 \times 256$
- Random crops of size  $227 \times 227$  were generated from inside the  $256 \times 256$  images to feed the first layer of AlexNet.
- Input:  $227 \times 227 \times 3$



# AlexNet



- 5 Convolutional Layers and 3 Fully Connected Layers
- It has 60 million parameters and 650,000 neurons and took five to six days to train on two GTX 580 3GB GPUs!
- Input:  $227 \times 227 \times 3$
- The first Conv Layer of AlexNet contains 96 kernels of size  $11 \times 11 \times 3$ .
  - Width and height of output:  $(227-11)/4 + 1 = 55$
- Number of parameters in first layer?
  - $(11 \times 11 \times 3 + 1) \times 96 = 34,944$
- Number of Computations:  $34,944 \times 55 \times 55 = 10,57,05,600$
- <https://airtable.com/shrArXKRCau4KhAwZ/tbloN5WYjFYpKGUEt?blocks=hide>

# References



1. Chapter 9, Deep Learning: Ian Goodfellow, Yoshua Bengio, Aaron Courville (<http://www.deeplearningbook.org/>)
2. <https://neurdivness.wordpress.com/2018/05/17/deep-convolutional-neural-networks-as-models-of-the-visual-system-qa/>
3. <https://www.rctn.org/bruno/public/papers/Fukushima1980.pdf>
4. <https://www.youtube.com/watch?v=2-Ol7ZBoMmU>
5. <https://www.youtube.com/watch?v=ZOXOwYUVCqw>

